

Machine Learning

Tusha Karnani

Table of contents

K-means clustering	1
Hierarchichal Clustering	6
Principal Component Analysis	7
PCA of UK food data	7
PCA	10


Fundamentals of machine learning including clustering and dimensionality reduction

K-means clustering

Creating a data set:

```
hist(rnorm(500,mean=5))
```

Histogram of rnorm(500, mean = 5)



```
x <- c(rnorm(30,mean=-3),rnorm(30,mean=3))
y <- rev(x)

data <- cbind(x,y)
```

```
k <- kmeans(data, 2)  
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```

          x           y
1 -3.011475  2.962974
2  2.962974 -3.011475

```

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 48.30937 48.30937
```

(between_SS / total_SS = 91.7 %)

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

How many points are in each cluster?

k\$size

[1] 30 30

What components of the result object details: - cluster assignment/membership? - cluster center?


k\$cluster

```
k$centers
```


	x	y
1	-3.011475	2.962974
2	2.962974	-3.011475

Clustering results figure colored by cluster membership looks like the following:

```
plot(data, col=c("red", "blue"))
```



```
plot(data, col=k$cluster)
points(k$centers, col="royalblue", pch=20, cex=2)
```




```
# pch is plotting character i.e. dot type
# cex is character expansion
```

K-means clustering is very popular as it is very fast and relatively straightforward: it takes numeric data as input and return the cluster membership, etc.

The problem is that we must tell k-means how many clusters we want.

Clustering it into 4 groups and plotting the results like above looks like the following:

```
k4 <- kmeans(data, 4)
plot(data, col=k4$cluster)
points(k4$centers, col="royalblue", pch=20, cex=2)
```




It picks lowest value of total variation within clusters. That is the tightest fit and returns that. There is a large reduction in SS at k=2 but after that the values do not go down as quickly.

```
k1 <- kmeans(data, 1)
k3 <- kmeans(data, 3)
k5 <- kmeans(data, 5)

ss <- c(k1$tot.withinss, k$tot.withinss, k3$tot.withinss, k4$tot.withinss, k5$tot.withinss)
ss
```

```
[1] 1167.43976    96.61873    74.10168    65.73590    61.03211
```


```
plot(ss, type="b", xlab="k", ylab="total within ss")
```



```
n <- NULL

for (i in 1:5)
{
  n <- c(n, kmeans(data, i)$tot.withinss)
}

plot(n, type="b")
```



Hierarchichal Clustering

We can't just input our data for this directly. We need to first calculate a distance matrix for our data (e.g. `dist()`) and use it as input for the `hclust()` function. Note: We can do this using not just euclidean distance but also sequence similarity, amino acid similarity, etc.

```
d <- dist(data)
# 60x60 matrix with the distance of every point in the dataset with every other point


hc <- hclust(d)
hc
```

Call:
`hclust(d = d)`

Cluster method : complete
Distance : euclidean
Number of objects: 60

```
plot(hc)
abline(h=8, col="red")
```

Cluster Dendrogram



d
`hclust (*, "complete")`


The numbers between 1-30 and 31-60 are on separate sides of the graphs i.e. they are highly separated by distance.

To get our cluster membership vector (i.e. our main clustering result) we can “cut” the tree at a given height or a height that yields a given “k” number of groups.

```
hccol <- cutree(hc, h=8)
```

Plot the data with our hclust result coloring

```
plot(data, col=hccol)
```



Principal Component Analysis

PCA of UK food data

For this, I imported a data set containing consumption patterns of 6 food groups across 4 regions.

```
url <- "https://tinyurl.com/UK-foods"
food <- read.csv(url, row.names=1)
dim(food)
```

```
[1] 17  4
```

```
head(food)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139


```
rownames(food) <- food [,1]
food <- food[,-1]
food
```

	Wales	Scotland	N.Ireland
105	103	103	66
245	227	242	267
685	803	750	586
147	160	122	93
193	235	184	209
156	175	147	139
720	874	566	1033
253	265	171	143
488	570	418	355
198	203	220	187
360	365	337	334
1102	1137	957	674
1472	1582	1462	1494
57	73	53	47
1374	1256	1572	1506
375	475	458	135
54	64	62	41


```
# Can get disruptive if run again (keeps removing columns)
```

Some base figures

```
barplot(as.matrix(food), beside=T, col=rainbow(nrow(food)))
```




```
barplot(as.matrix(food), beside=F, col=rainbow(nrow(food)))
```



There is one plot that can be useful for small datasets

```
pairs(food, col=rainbow(10), pch=16)
```



Note: It can be challenging to spot major trends and patterns even in relatively small multi-variate datasets (here we have 17, typically we have 1000s)

PCA

In base R, the main PCA function is `prcomp()`

I will take the transpose of our food data so the “foods” are in the columns

```
pca <- prcomp(t(food))
summary(pca)
```


Importance of components:

	PC1	PC2	PC3
Standard deviation	379.8991	260.5533	1.438e-13
Proportion of Variance	0.6801	0.3199	0.000e+00
Cumulative Proportion	0.6801	1.0000	1.000e+00

```
pca$x
```


	PC1	PC2	PC3
Wales	-288.9534	226.36855	1.776357e-15
Scotland	-141.3603	-284.81172	4.298784e-13
N.Ireland	430.3137	58.44317	-9.592327e-14

```
cols <- c("black","red","blue")
plot(pca$x[,1], pca$x[,2], col=cols, pch=16, xlab="PCA1", ylab="PCA2")
```



```
library(ggplot2)
```

```
ggplot(pca$x) +
  aes(PC1, PC2) +
  geom_point(col=cols) +
  theme_minimal()
```



```
pca$rotation
```

	PC1	PC2	PC3
105	-0.05515951	-0.015926208	0.089159718
245	0.05228588	-0.014247351	0.160169691
685	-0.29754711	0.017770827	-0.143714947
147	-0.08127351	0.050871661	-0.035378388
193	-0.01378426	0.095789177	-0.088149169
156	-0.03995614	0.043238659	0.062522585
720	0.38787509	0.714518360	-0.188185663
253	-0.13584223	0.144666425	-0.044940419
488	-0.24608178	0.226299860	-0.277334193
198	-0.03217825	-0.042547197	-0.092922896
360	-0.03250215	0.045390849	0.144304277
1102	-0.60208698	0.178285653	0.102836115
1472	-0.07242201	0.213840430	-0.659763755
57	-0.02896604	0.030761774	0.009935304
1374	0.21794320	-0.555250465	-0.557926740
375	-0.49854565	-0.110688746	-0.178079520
54	-0.03330887	-0.005704759	0.006954684

```
# How much each of these variables weighs on each of the axes/pcs  
# Telling us what makes two countries different from each other (based on PCA results)
```

```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col() +  
  theme_minimal()
```

